


Four fast models, one harder routing test

500 English messages • 77 banking intents • real API calls



Latency includes network and schema validation. One outstanding call per provider, four providers in parallel.

Costs use reported charges where available; otherwise usage estimates or error reservations. Timing repeats and development excluded.

Failed calls count as incorrect. Latency is for valid answers only. One dataset, one host, one run.