



Do stated probabilities match observed accuracy?

Five fixed probability bins. Bars show Wilson 95% intervals, not model uncertainty.

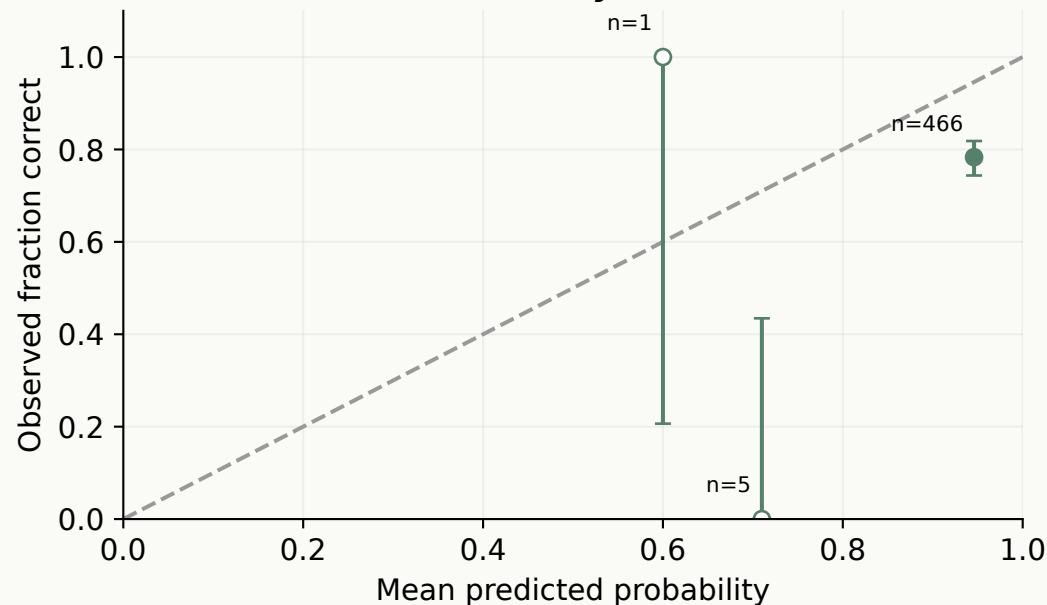
Jev 1.13



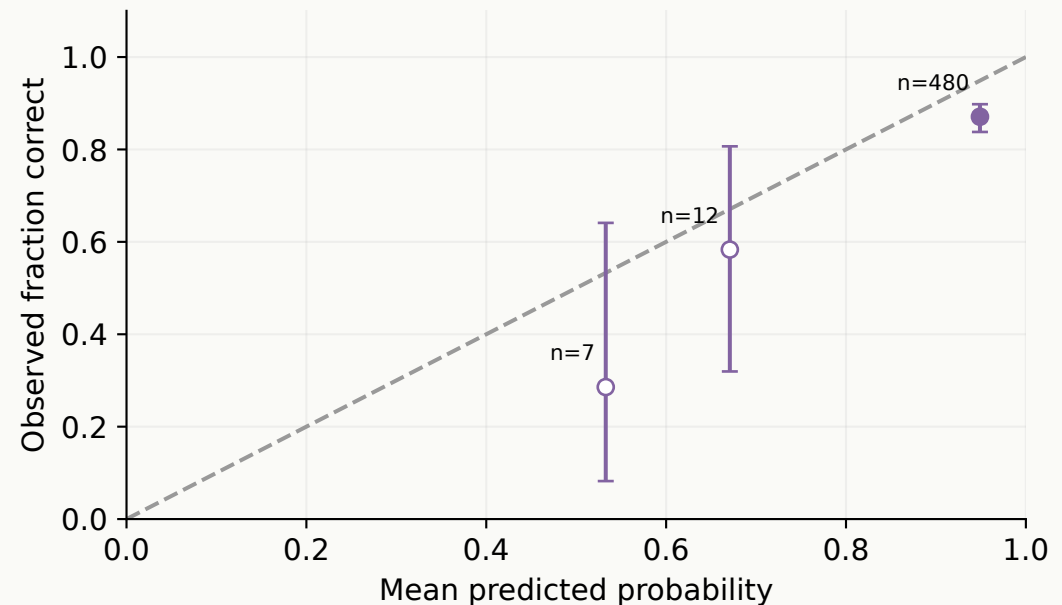
GPT-OSS / Cerebras



Mercury 2.5



Gemini 3.8 Flash



Hollow points: fewer than 20 examples. LLMs report verbalized probabilities; Jev supplies class probabilities.

Sparse bins and public-dataset contamination limit conclusions. Native Jev confidence is not used in this plot.