

The test, in order

BANKING77 · English only · four pinned model/provider pairs

50

Development

Check requests, parsing, confidence capture, and budget logic.



200

Choose cutoffs

Separate training-split messages. Freeze the acceptance rules.



500

Held-out test

Official test-split subset. Do not change the frozen cutoffs.



50

Timing repeats

Repeat a fixed subset. Not additional accuracy evidence.

Counts are per model. 3,200 API attempts total; 500 independent test messages.
No retries or fallback. No AI-generated answer key for this follow-up.