


How often did each system get it right?

One attempt per system on the same 500 held-out messages.



Bars: accuracy over all 500 attempts. Lines: 95% paired, intent-stratified bootstrap intervals.

Failures count as incorrect: Jev 0, GPT-OSS 0, Mercury 28, Gemini 1.