



What did the frozen rules accept?

Cutoffs selected on 200 separate messages; evaluated on 500 unseen messages.

Jev uses native confidence. Gemini uses its reported probability.



Both intervals extend above the 5% target (dashed line). Neither rule establishes a risk guarantee.

GPT-OSS and Mercury found no qualifying cutoff on calibration; this does not mean their scores cannot rank errors.